

ZKPrivacy

AI Implementation & Verification Guide

Version 1.0

Draft – January 2026

Phexora AI

<https://quantum.phexora.ai>

Companion document to the ZKPrivacy Technical Specification.
Designed for AI-assisted implementation and verification.

Contents

1	ZKPrivacy: AI Implementation & Verification Guide	1
1.1	Purpose	1
2	CRITICAL: IMMUTABLE REQUIREMENTS CHECK	1
2.1	BEFORE ANY IMPLEMENTATION	1
2.1.1	Pre-Implementation Checklist	1
2.1.2	Automatic Requirement Violation Detection	2
2.1.3	Requirement Violation = Implementation Failure	3
3	Part I: Implementation Task Decomposition	3
3.1	1. Dependency Graph	3
3.2	2. Task Specifications	5
3.2.1	T001: Field Arithmetic (Goldilocks)	5
3.2.2	T002: Polynomial Ring Arithmetic	5
3.2.3	T003: Hash Functions	6
3.2.4	T101: NTT Implementation	7
3.2.5	T102: Merkle Tree	7
3.2.6	T201: Lattice Commitment Scheme	8
3.2.7	T301: Output Structure	9
3.2.8	T401: Transaction Structure	10
3.2.9	T501: STARK Prover	11
3.2.10	T502: STARK Verifier	12
3.3	3. Integration Tests	12
3.3.1	IT001: End-to-End Transaction	12
3.3.2	IT002: Double-Spend Prevention	12
3.3.3	IT003: Chain Reorganization	13
4	Part II: Verification Oracles	13
4.1	4. Cryptographic Oracles	13
4.1.1	O001: Commitment Binding Oracle	13
4.1.2	O002: STARK Soundness Oracle	14
4.1.3	O003: Privacy Oracle	14
4.2	5. Performance Oracles	15
4.2.1	O004: Proof Generation Benchmark	15
4.2.2	O005: Verification Benchmark	16
4.3	6. Fuzz Testing Specifications	17
4.3.1	F001: Serialization Fuzzing	17
4.3.2	F002: Arithmetic Fuzzing	17
5	Part III: Quality Metrics	18

5.1	7. Code Quality Checklist	18
5.2	8. Security Checklist	18
5.3	9. Performance Metrics	19
6	Part IV: Acceptance Criteria	19
6.1	10. Minimum Viable Implementation	19
6.2	11. Formal Verification Requirements	20
6.3	12. Benchmarking Protocol	21
6.4	13. Self-Verification Commands	22
7	End of Verification Guide	23

1 ZKPrivacy: AI Implementation & Verification Guide

1.1 Purpose

This document provides:

1. Structured implementation tasks for AI systems
2. Verification criteria and test oracles
3. Self-check mechanisms for implementation correctness
4. Benchmark metrics for quality assessment

Target: Advanced AI systems capable of autonomous code generation and verification.

2 CRITICAL: IMMUTABLE REQUIREMENTS CHECK

2.1 ☐ BEFORE ANY IMPLEMENTATION ☐

The implementing AI system MUST first verify understanding of the immutable requirements (Section R1-R7 of the main specification). These requirements are **non-negotiable** and **cannot be modified**.

2.1.1 Pre-Implementation Checklist

Before writing ANY code, the AI MUST confirm:

[] R1: PRIVACY REQUIREMENTS

- [] R1.1: No transparent transaction mode will be implemented
- [] R1.2: Anonymity set equals entire output set (not rings/subsets)
- [] R1.3: Outputs are unlinkable to addresses without view key
- [] R1.4: All amounts are hidden via commitments
- [] R1.5: Dandelion++ is included in network layer

[] R2: SECURITY REQUIREMENTS

- [] R2.1: ONLY post-quantum primitives used:
 - Commitments: Lattice-based (Module-LWE)
 - Signatures: SPHINCS+ (hash-based)
 - Key exchange: ML-KEM (Kyber)
 - ZK proofs: STARKs (hash-based)
 - NO elliptic curves anywhere
- [] R2.2: No trusted setup exists
- [] R2.3: Commitment binding is proven

- [] R2.4: Commitment hiding is proven
- [] R2.5: STARK soundness $< 2^{-100}$
- [] R2.6: STARK is zero-knowledge

- [] R3: DECENTRALIZATION REQUIREMENTS
 - [] R3.1: No permissions required to participate
 - [] R3.2: No privileged parties or special keys
 - [] R3.3: RandomX (ASIC-resistant) for mining
 - [] R3.4: All code is open source

- [] R4: INTEGRITY REQUIREMENTS
 - [] R4.1: Supply capped at exactly 21,000,000 ZKP
 - [] R4.2: Balance equation enforced in ZK proof
 - [] R4.3: Nullifiers prevent double-spending
 - [] R4.4: Heaviest chain rule for finality

- [] R5: FUNCTIONAL REQUIREMENTS
 - [] R5.1: Multi-input/output transactions supported
 - [] R5.2: Deterministic wallets from seed
 - [] R5.3: Light client support

- [] R6: PERFORMANCE REQUIREMENTS
 - [] R6.1: Proof generation $< 120s$, verification $< 2s$
 - [] R6.2: Runs on 16GB RAM, 500GB storage
 - [] R6.3: 120s block time, 10 TPS

- [] R7: NON-REQUIREMENTS ACKNOWLEDGED
 - [] R7.1: No smart contracts in v1
 - [] R7.2: No backdoors, no censorship, no blacklisting

2.1.2 Automatic Requirement Violation Detection

The implementation MUST include automated checks that detect requirement violations:

```
// Run at compile time and runtime
pub fn verify_requirements_compliance() -> Result<(), RequirementViolation> {
    // R1.1: No transparent mode
    #[cfg(feature = "transparent_transactions")]
    compile_error!("REQUIREMENT VIOLATION R1.1: Transparent transactions are forbidden");

    // R2.1: No elliptic curves
    #[cfg(feature = "elliptic_curves")]
    compile_error!("REQUIREMENT VIOLATION R2.1: Elliptic curves are forbidden");
}
```

```

#[cfg(feature = "elliptic_curves")]
compile_error!("REQUIREMENT VIOLATION R2.1: Elliptic curves are forbidden");

// R4.1: Supply cap
const MAX_SUPPLY: u64 = 21_000_000 * 100_000_000; // in satoshi
assert!(TOTAL_SUPPLY_CAP == MAX_SUPPLY, "REQUIREMENT VIOLATION R4.1");

// R3.2: No privileged keys
assert!(PRIVILEGED_KEYS.is_empty(), "REQUIREMENT VIOLATION R3.2");

Ok(())
}

```

2.1.3 Requirement Violation = Implementation Failure

IF any requirement R1-R7 is violated:
 THEN implementation is INVALID
 THEN implementation MUST be rejected
 THEN implementation MUST be restarted from scratch

There is NO partial compliance.

There are NO exceptions.

There are NO workarounds.

3 Part I: Implementation Task Decomposition

3.1 1. Dependency Graph

Implementation order (topological sort):

Level 0 (No dependencies):

- T001: Field arithmetic (Goldilocks)
- T002: Polynomial ring arithmetic
- T003: Hash functions (SHAKE256 wrappers)
- T004: Serialization primitives

Level 1 (Depends on Level 0):

- T101: NTT implementation

T102: Merkle tree
T103: Domain-separated hash instances
T104: Random number generation

Level 2 (Depends on Level 1):

T201: Lattice commitment scheme
T202: SPHINCS+ integration
T203: ML-KEM integration
T204: Nullifier derivation

Level 3 (Depends on Level 2):

T301: Output structure
T302: Stealth addresses
T303: Key hierarchy
T304: Address encoding

Level 4 (Depends on Level 3):

T401: Transaction structure
T402: STARK AIR definition
T403: Wallet scanning
T404: Transaction creation

Level 5 (Depends on Level 4):

T501: STARK prover
T502: STARK verifier
T503: Block structure
T504: Consensus rules

Level 6 (Depends on Level 5):

T601: Full node
T602: P2P networking
T603: Chain state management
T604: Complete wallet

Level 7 (Integration):

T701: System integration and testing

3.2 2. Task Specifications

3.2.1 T001: Field Arithmetic (Goldilocks)

Input specification:

Field: F_p where $p = 2^{64} - 2^{32} + 1$
Elements: 64-bit unsigned integers

Required operations:

```
trait GoldilocksField {  
    fn add(a: u64, b: u64) -> u64;  
    fn sub(a: u64, b: u64) -> u64;  
    fn mul(a: u64, b: u64) -> u64;  
    fn inv(a: u64) -> u64; // Multiplicative inverse  
    fn pow(base: u64, exp: u64) -> u64;  
    fn neg(a: u64) -> u64;  
  
    // Batch operations for efficiency  
    fn batch_inv(elements: &[u64]) -> Vec<u64>;  
}
```

Verification criteria:

V001.1: $\text{add}(a, b) = (a + b) \bmod p$
V001.2: $\text{mul}(a, b) = (a \times b) \bmod p$
V001.3: $\text{inv}(a) \times a = 1 \bmod p$ for $a \neq 0$
V001.4: All operations complete in constant time
V001.5: No overflow in intermediate computations

Test vectors:

```
add(p-1, 1) = 0  
mul(p-1, p-1) = 1  
inv(2) = (p+1)/2 = 9223372034707292161  
pow(7, p-1) = 1 (Fermat's little theorem)
```

3.2.2 T002: Polynomial Ring Arithmetic

Input specification:

Ring: $R_q = \mathbb{Z}_q[X]/(X^{256} + 1)$
 $q = 8380417$
Coefficients: Signed 32-bit integers (reduced mod q)

Required operations:

```
trait PolynomialRing {
    fn add(a: &Poly, b: &Poly) -> Poly;
    fn sub(a: &Poly, b: &Poly) -> Poly;
    fn mul(a: &Poly, b: &Poly) -> Poly; // Via NTT
    fn scalar_mul(a: &Poly, s: i32) -> Poly;
    fn reduce(a: &Poly) -> Poly; // Reduce coefficients mod q
}

type Poly = [i32; 256];
```

Verification criteria:

- V002.1: Coefficients always in $[-(q-1)/2, (q-1)/2]$ after reduce
 - V002.2: mul satisfies $(X^{256} + 1)$ reduction
 - V002.3: Ring axioms hold (associativity, distributivity)
-

3.2.3 T003: Hash Functions

Input specification:

Base: SHAKE256 (FIPS 202)

Domain separation: Prefix with tagged length-encoded domain string

Required instances:

```
trait HashInstances {
    fn h_commitment(input: &[u8]) -> [u8; 64];
    fn h_nullifier(input: &[u8]) -> [u8; 32];
    fn h_merkle(input: &[u8]) -> [u8; 32];
    fn h_address(input: &[u8]) -> [u8; 32];
    fn h_kdf(input: &[u8], output_len: usize) -> Vec<u8>;
    fn h_challenge(input: &[u8]) -> [u8; 64];
    fn h_pow(input: &[u8]) -> [u8; 32];
}
```

Verification criteria:

- V003.1: $h_X(m) \neq h_Y(m)$ for $X \neq Y$ (domain separation)
- V003.2: Output matches SHAKE256 reference implementation
- V003.3: Streaming API for large inputs

Test vectors (must match specification Section 18.1):

h_nullifier(0x00⁶⁴):

```

Domain: "ZKPrivacy-v1.nullifier"
Output: 0x3a7f2c9e8b4d1a6f5c0e7b3d9a2f8c4e1b6d0a5f3e9c7b2d8a4e6f1c0b5d9a3e

h_merkle(0x00 || 0x00^32 || 0x00^32):
    Domain: "ZKPrivacy-v1.merkle"
    Output: 0x5c9a3e7f1b4d8c2e6a0f5b9d3c7e1a4f8b2d6e0a4c9f3b7e1d5a8c2f6e0b4d9a

Domain separation test:
    h_nullifier(0x00^32)    h_commitment(0x00^32)    h_merkle(0x00^32)

```

3.2.4 T101: NTT Implementation

Input specification:

Transform size: 256
Field: Z_q where q = 8380417
Primitive 512th root of unity: = 1753

Required operations:

```

trait NTT {
    fn forward(a: &Poly) -> Poly;
    fn inverse(a: &Poly) -> Poly;
    fn pointwise_mul(a: &Poly, b: &Poly) -> Poly; // In NTT domain
}

```

Verification criteria:

V101.1: inverse(forward(a)) = a
V101.2: forward(a * b) = pointwise_mul(forward(a), forward(b))
V101.3: ^512 = 1 mod q
V101.4: ^256 = -1 mod q

3.2.5 T102: Merkle Tree

Input specification:

Hash: h_merkle
Depth: 40
Leaf prefix: 0x01
Node prefix: 0x00

Required operations:

```

trait MerkleTree {
    fn new(depth: u32) -> Self;
    fn append(&mut self, leaf: &[u8; 32]) -> u64; // Returns position
    fn root(&self) -> [u8; 32];
    fn prove(&self, position: u64) -> MerkleProof;
    fn verify(root: &[u8; 32], leaf: &[u8; 32], position: u64, proof: &MerkleProof) -> bool
}

struct MerkleProof {
    siblings: Vec<[u8; 32]>, // Length = depth
    path_bits: u64,           // Left/right indicators
}

```

Verification criteria:

- V102.1: verify(tree.root(), leaf, pos, tree.prove(pos)) = true
 - V102.2: Proof size exactly depth × 32 bytes
 - V102.3: Empty tree root is well-defined
 - V102.4: Append-only (no modification of existing leaves)
-

3.2.6 T201: Lattice Commitment Scheme

Input specification:

Parameters from spec: n=256, q=8380417, k=4, =2
 Public parameters: Matrix A R_q^{k×k}

Required operations:

```

trait LatticeCommitment {
    fn setup(seed: &[u8; 32]) -> PublicParams;
    fn commit(pp: &PublicParams, value: u64, randomness: &PolyVec) -> Commitment;
    fn verify_opening(pp: &PublicParams, c: &Commitment, v: u64, r: &PolyVec) -> bool
    fn add(c1: &Commitment, c2: &Commitment) -> Commitment;

    fn generate_randomness(seed: &[u8; 32]) -> PolyVec;
}

type PolyVec = [Poly; 4];
struct Commitment { data: PolyVec }

```

Verification criteria:

- V201.1: Randomness coefficients in [- ,]

V201.2: $\text{commit}(v_1, r_1) + \text{commit}(v_2, r_2) = \text{commit}(v_1+v_2, r_1+r_2)$
V201.3: Cannot find collision (binding)
V201.4: Commitment reveals nothing about value (hiding)

3.2.7 T301: Output Structure

Required structure:

```
struct Output {
    commitment: Commitment,           // ~3KB
    kyber_ciphertext: [u8; 1568],     // ML-KEM-1024
    encrypted_data: [u8; 128],         // AES-GCM
}

struct OutputPlaintext {
    value: u64,
    blinding_seed: [u8; 32],
    memo: [u8; 64],
    checksum: [u8; 16],
}

impl Output {
    fn create(
        pp: &PublicParams,
        recipient_view_pk: &KyberPublicKey,
        value: u64,
        memo: &[u8; 64],
    ) -> (Self, OutputSecrets);

    fn try_decrypt(
        &self,
        view_sk: &KyberSecretKey,
    ) -> Option<OutputPlaintext>;

    fn serialize(&self) -> Vec<u8>;
    fn deserialize(data: &[u8]) -> Result<Self, Error>;
}
```

Verification criteria:

V301.1: Serialization is canonical and deterministic

V301.2: try_decrypt succeeds only with correct key
V301.3: Checksum validates integrity
V301.4: Output size matches specification (~13KB)

3.2.8 T401: Transaction Structure

Required structure:

```
struct Transaction {
    nullifiers: Vec<[u8; 32]>,
    outputs: Vec<Output>,
    fee: u64,
    validity_proof: StarkProof,
    authorization: Authorization,
    anchor: [u8; 32],
}

impl Transaction {
    fn create(
        wallet: &Wallet,
        inputs: &[OwnedOutput],
        recipients: &[(Address, u64)],
        fee: u64,
        anchor: [u8; 32],
    ) -> Result<Self, Error>;
    fn verify(&self, utxo_tree_root: &[u8; 32]) -> bool;
    fn nullifier_count(&self) -> usize;
    fn output_count(&self) -> usize;
    fn serialized_size(&self) -> usize;
}
```

Verification criteria:

V401.1: Valid transaction passes verify()
V401.2: Invalid balance fails verify()
V401.3: Reused nullifier fails verify()
V401.4: Wrong anchor fails verify()

3.2.9 T501: STARK Prover

Input specification:

Statement: Transaction validity (balance, range, membership, nullifiers)

Security: 100-bit soundness

Proof size: < 200 KB

Required interface:

```
trait StarkProver {
    fn prove(
        public_inputs: &PublicInputs,
        witness: &Witness,
    ) -> StarkProof;
}

struct PublicInputs {
    nullifiers: Vec<[u8; 32]>,
    output_commitments: Vec<Commitment>,
    fee: u64,
    anchor: [u8; 32],
}

struct Witness {
    input_values: Vec<u64>,
    input_bindings: Vec<PolyVec>,
    input_positions: Vec<u64>,
    input_merkle_paths: Vec<MerkleProof>,
    input_nullifier_keys: Vec<[u8; 32]>,
    output_values: Vec<u64>,
    output_bindings: Vec<PolyVec>,
}
```

Verification criteria:

V501.1: Proof verifies for valid witness

V501.2: Cannot generate valid proof for invalid statement

V501.3: Proof size < 200 KB

V501.4: Proving time < 60 seconds (benchmark hardware)

3.2.10 T502: STARK Verifier

Required interface:

```
trait StarkVerifier {  
    fn verify(  
        public_inputs: &PublicInputs,  
        proof: &StarkProof,  
    ) -> bool;  
}
```

Verification criteria:

V502.1: Accepts valid proofs

V502.2: Rejects invalid proofs with overwhelming probability

V502.3: Verification time < 1 second

V502.4: Constant-time execution (no timing leaks)

3.3 3. Integration Tests

3.3.1 IT001: End-to-End Transaction

Setup:

1. Generate two wallets (Alice, Bob)
2. Initialize chain with genesis block
3. Mine blocks to give Alice coins

Test:

4. Alice creates transaction sending to Bob
5. Transaction is validated and included in block
6. Bob's wallet scans and finds output
7. Bob can spend the received output

Verify:

- Alice's balance decreased correctly
- Bob's balance increased correctly
- Nullifiers are recorded
- Chain state is consistent

3.3.2 IT002: Double-Spend Prevention

Setup:

1. Wallet with single UTXO

Test:

2. Create transaction spending the UTXO
3. Create second transaction spending same UTXO
4. Submit both to node

Verify:

- First transaction accepted
- Second transaction rejected
- Only one nullifier recorded

3.3.3 IT003: Chain Reorganization

Setup:

1. Chain at height 100
2. Fork at height 95 with more cumulative work

Test:

3. Receive fork blocks
4. Process reorganization

Verify:

- Chain switches to fork
 - Transactions in orphaned blocks return to mempool
 - State correctly reflects new chain
-

4 Part II: Verification Oracles

4.1 4. Cryptographic Oracles

4.1.1 O001: Commitment Binding Oracle

```
def test_binding(implementation):  
    """Test that commitments are binding."""  
    pp = implementation.setup(random_seed())  
  
    # Try to find collision  
    for _ in range(1000000):  
        v1, r1 = random_value(), random_blinding()  
        v2, r2 = random_value(), random_blinding()
```

```

c1 = implementation.commit(pp, v1, r1)
c2 = implementation.commit(pp, v2, r2)

if c1 == c2 and (v1, r1) != (v2, r2):
    return FAIL("Found collision")

return PASS("No collision found in 10^6 attempts")

```

4.1.2 O002: STARK Soundness Oracle

```

def test_soundness(implementation):
    """Test STARK soundness with invalid witnesses."""

    test_cases = [
        # Unbalanced transaction
        {"inputs": [100], "outputs": [50, 60], "fee": 0},  # 110 > 100

        # Negative value (overflow attempt)
        {"inputs": [100], "outputs": [2**64-1, 101], "fee": 0},

        # Invalid Merkle proof
        {"valid_merkle": False},
    ]

    for case in test_cases:
        witness = generate_invalid_witness(case)
        proof = implementation.prove(witness)

        if implementation.verify(proof):
            return FAIL(f"Accepted invalid case: {case}")

    return PASS("Rejected all invalid cases")

```

4.1.3 O003: Privacy Oracle

```

def test_transaction_privacy(implementation):
    """Test that transactions reveal no private information."""

```

```

# Create many transactions with different parameters
transactions = []
for _ in range(1000):
    tx = implementation.create_transaction(
        random_inputs(),
        random_outputs(),
        random_fee()
    )
    transactions.append(tx)

# Statistical tests on serialized transactions
serialized = [tx.serialize() for tx in transactions]

# Test 1: No correlation between tx size and value
if correlation(sizes, values) > 0.1:
    return FAIL("Size leaks value information")

# Test 2: Byte distribution is uniform
if not chi_squared_uniform(concatenate(serialized)):
    return FAIL("Non-uniform byte distribution")

# Test 3: No timing correlation
prove_times = measure_prove_times(transactions)
if correlation(prove_times, values) > 0.1:
    return FAIL("Timing leaks value information")

return PASS("No detectable information leakage")

```

4.2 5. Performance Oracles

4.2.1 O004: Proof Generation Benchmark

```

def benchmark_proving(implementation):
    """Benchmark proof generation time."""

configurations = [
    {"inputs": 1, "outputs": 2},    # Minimal
    {"inputs": 2, "outputs": 2},    # Typical
    {"inputs": 4, "outputs": 2},    # Multi-input

```

```

        {"inputs": 2, "outputs": 8},    # Multi-output
        {"inputs": 16, "outputs": 2},   # Large
    ]

results = {}
for config in configurations:
    times = []
    for _ in range(10):
        witness = generate_valid_witness(config)
        start = time.monotonic()
        implementation.prove(witness)
        elapsed = time.monotonic() - start
        times.append(elapsed)

    results[str(config)] = {
        "mean": statistics.mean(times),
        "std": statistics.stdev(times),
        "max": max(times),
    }

# Check against requirements
for config, result in results.items():
    if result["mean"] > 60.0:
        return FAIL(f"Proving too slow for {config}: {result['mean']}s")

return PASS(f"All configurations within limits: {results}")

```

4.2.2 O005: Verification Benchmark

```

def benchmark_verification(implementation):
    """Benchmark proof verification time."""

proofs = [generate_valid_proof() for _ in range(100)]

times = []
for proof in proofs:
    start = time.monotonic()
    result = implementation.verify(proof)
    elapsed = time.monotonic() - start
    times.append(elapsed)
    assert result == True

```

```

mean_time = statistics.mean(times)
max_time = max(times)

if max_time > 1.0:
    return FAIL(f"Verification too slow: max {max_time:.3f}s")

return PASS(f"Verification time: mean={mean_time:.3f}s, max={max_time:.3f}s")

```

4.3 6. Fuzz Testing Specifications

4.3.1 F001: Serialization Fuzzing

```

def fuzz_serialization(implementation):
    """Fuzz test serialization/deserialization."""

    # Property: deserialize(serialize(x)) == x
    for _ in range(1000000):
        tx = generate_random_transaction()
        serialized = implementation.serialize(tx)
        deserialized = implementation.deserialize(serialized)
        assert tx == deserialized

    # Property: Invalid bytes should not crash
    for _ in range(1000000):
        random_bytes = os.urandom(random.randint(0, 1000000))
        try:
            implementation.deserialize(random_bytes)
        except DeserializationError:
            pass # Expected
        except Exception as e:
            return FAIL(f"Unexpected exception: {e}")

    return PASS("Serialization fuzzing passed")

```

4.3.2 F002: Arithmetic Fuzzing

```

def fuzz_field_arithmetic(implementation):
    """Fuzz test field operations."""

```

```

p = 2**64 - 2**32 + 1

for _ in range(10000000):
    a = random.randint(0, p-1)
    b = random.randint(0, p-1)

    # Addition
    assert implementation.add(a, b) == (a + b) % p

    # Multiplication
    assert implementation.mul(a, b) == (a * b) % p

    # Inverse
    if a != 0:
        inv_a = implementation.inv(a)
        assert implementation.mul(a, inv_a) == 1

return PASS("Field arithmetic fuzzing passed")

```

5 Part III: Quality Metrics

5.1 7. Code Quality Checklist

- [] CQ001: No compiler warnings (strict mode)
- [] CQ002: All public APIs documented
- [] CQ003: Test coverage > 80%
- [] CQ004: No unsafe blocks (or justified and audited)
- [] CQ005: No panics in library code
- [] CQ006: All errors are typed and recoverable
- [] CQ007: No hardcoded magic numbers (use named constants)
- [] CQ008: Consistent naming conventions
- [] CQ009: No dead code
- [] CQ010: Passes clippy/lint with no warnings

5.2 8. Security Checklist

- [] SC001: All secret operations are constant-time
- [] SC002: Sensitive memory is zeroized after use
- [] SC003: No secret-dependent branches

- [] SC004: No secret-dependent memory access patterns
- [] SC005: CSPRNG used for all randomness
- [] SC006: Input validation on all external data
- [] SC007: No integer overflow vulnerabilities
- [] SC008: No buffer overflow vulnerabilities
- [] SC009: Rate limiting on network inputs
- [] SC010: Timeout handling on all operations

5.3 9. Performance Metrics

Target metrics (consumer hardware: 8-core CPU, 16GB RAM):

Operation	Target	Measured
Field multiplication	< 10 ns	
Polynomial multiplication	< 100 s	
Commitment creation	< 10 ms	
Merkle proof generation	< 1 ms	
Merkle proof verification	< 1 ms	
Output encryption	< 5 ms	
Output decryption	< 5 ms	
Wallet scan (per output)	< 5 ms	
STARK proof generation (2-2)	< 60 s	
STARK proof verification	< 1 s	
Block validation (1000 tx)	< 10 s	
Transaction serialization	< 1 ms	
Transaction deserialization	< 1 ms	

6 Part IV: Acceptance Criteria

6.1 10. Minimum Viable Implementation

An implementation is considered **complete** when:

Cryptographic Layer:

- [x] All hash functions implemented with test vectors passing
- [x] Lattice commitment scheme with binding/hiding tests
- [x] ML-KEM integration with KAT vectors
- [x] SPHINCS+ integration with KAT vectors
- [x] Merkle tree with correctness tests

Transaction Layer:

- [x] Output creation and encryption
- [x] Output scanning and decryption
- [x] Nullifier derivation
- [x] Transaction structure serialization

Proof System:

- [x] AIR constraints correctly specified
- [x] STARK prover generates valid proofs
- [x] STARK verifier rejects invalid proofs
- [x] Soundness tests pass

Consensus:

- [x] Block header validation
- [x] Difficulty adjustment
- [x] Chain selection rules
- [x] Block validation with transaction verification

Networking:

- [x] Peer discovery
- [x] Block propagation
- [x] Transaction propagation
- [x] Dandelion++ for transaction privacy

Wallet:

- [x] Key generation from seed
- [x] Address derivation
- [x] Balance tracking
- [x] Transaction creation
- [x] Output scanning

Integration:

- [x] End-to-end transaction test passes
- [x] Chain synchronization works
- [x] No memory leaks in long-running tests
- [x] All performance targets met

6.2 11. Formal Verification Requirements

Mandatory formal verification targets:

1. Field arithmetic correctness
Tool: Bounded model checking (Kani)
Property: All operations produce correct results mod p
2. Memory safety
Tool: MIRI
Property: No undefined behavior
3. Merkle tree correctness
Tool: Property-based testing (QuickCheck/Proptest)
Property: Verify-proof always succeeds for valid proofs
4. STARK verifier soundness
Tool: Cryptographic review + testing
Property: Rejects invalid proofs
5. No secret-dependent timing
Tool: dudect or similar
Property: Timing is independent of secret values

Optional formal verification targets:

6. Full STARK soundness proof (EasyCrypt/Coq)
 7. Protocol-level security proof (TLA+)
 8. Economic incentive analysis (game theory)
-

6.3 12. Benchmarking Protocol

For reproducible benchmarks:

Hardware specification:

CPU: AMD Ryzen 9 5900X or equivalent
RAM: 64 GB DDR4-3200
Storage: NVMe SSD
OS: Ubuntu 22.04 LTS

Software configuration:

Rust: Latest stable
Compiler flags: --release with LTO
No other significant processes running

Benchmark procedure:

1. Warm-up: Run operation 100 times, discard results
 2. Measurement: Run operation 1000 times
 3. Report: Mean, median, std dev, min, max
 4. Verify: Results reproducible across runs
-

6.4 13. Self-Verification Commands

An AI implementation should expose these verification commands:

```
# Run all unit tests
zkprivacy test --all

# Run cryptographic test vectors
zkprivacy test --vectors

# Run fuzz tests (1 hour)
zkprivacy fuzz --duration 3600

# Run performance benchmarks
zkprivacy bench --full

# Run security checks
zkprivacy audit --security

# Verify against specification
zkprivacy verify --spec ./zkprivacy-quantum-spec-v1.md

# Generate implementation report
zkprivacy report --output implementation-report.json
```

Expected report format:

```
{  
    "version": "1.0.0",  
    "spec_version": "1.0",  
    "timestamp": "2026-01-14T12:00:00Z",  
    "tests": {  
        "unit": {"passed": 1234, "failed": 0, "skipped": 0},  
        "integration": {"passed": 56, "failed": 0, "skipped": 0},  
        "fuzz": {"iterations": 10000000, "failures": 0}  
    }  
}
```

```
},
"benchmarks": {
    "proof_generation_2_2": {"mean_ms": 45000, "target_ms": 60000},
    "proof_verification": {"mean_ms": 800, "target_ms": 1000}
},
"security": {
    "constant_time_check": "PASS",
    "memory_safety": "PASS",
    "input_validation": "PASS"
},
"coverage": {
    "line": 0.87,
    "branch": 0.82
},
"verification_hash": "sha256:abc123..."
}
```

7 End of Verification Guide

This document provides complete criteria for implementing and verifying the ZKPrivacy specification. An implementation that satisfies all requirements in this document is considered conformant.