

Quantum Research Implementation and Verification Plan

Evidence plan for the 0.4.0 research design

Phexora AI · phexora.ai

2026-07-23

Abstract

This document turns the Quantum research requirements into a human-owned, dependency-ordered implementation and verification plan. It defines reviewer roles, evidence classes, task gates, adversarial scenarios, performance reporting, and release criteria for a candidate private post-quantum DAG protocol. It is a planning and assurance document; it does not report a conformant implementation or completed verification.

Keywords: post-quantum cryptography; protocol verification; evidence gates; reproducible benchmarking; independent review.

1. Purpose and status

This guide turns the Quantum research requirements into an ordered body of work. It does **not** certify the design, prove security, or report completed tests. At this revision:

- there is no protocol implementation in this repository;
- none of the task acceptance gates has passed;
- no benchmark has been run;
- no external cryptographic or consensus review has been completed;
- there is no conformant testnet or production network.

The goal is a protocol that simultaneously satisfies:

1. end-to-end post-quantum security;
2. private ledger semantics and network anonymity by default; and
3. at least 1,000 accepted layer-1 transactions per second.

These are hard release requirements. They are not current capability claims. The initial product boundary is private post-quantum cash and settlement with bounded payment policies, not a general-purpose smart-contract platform.

2. Governance and reviewer roles

AI systems may help draft code, tests, vectors, and analyses. They **MUST NOT** approve their own cryptographic design, formal argument, benchmark, audit finding, or release.

Every task has:

- an accountable human owner;
- an independent reviewer who did not author the work;
- explicit dependencies;
- versioned artifacts;
- a pass/fail decision with evidence;
- unresolved-risk and change-control records.

The minimum specialist review roles are:

Role	Approval boundary
Cryptography reviewer	primitives, parameters, KATs, composition
Proof-system reviewer	AIR, transcript, FRI, ZK, soundness

Role	Approval boundary
Consensus reviewer	GHOSTDAG profile, DAA, ordering, finality
Network-privacy reviewer	observer model, routing, traffic analysis
Implementation-security reviewer	parsers, side channels, supply chain
Performance reviewer	workload, topology, measurement validity
Legal counsel	launch, token, privacy, sanctions and jurisdiction questions

No single reviewer may sign all technical boundaries.

3. Evidence taxonomy

Reports MUST label evidence accurately:

Evidence type	What it can support	What it cannot support alone
Analytical proof	a stated theorem under stated assumptions	implementation correctness
Formal verification	properties encoded and proved in a named formal tool	omitted properties or real hardware behavior
KAT/conformance vector	exact algorithm and encoding agreement	general security
Unit/property test	exercised implementation behavior	negligible failure probability or a reduction
Differential/fuzz test	cross-implementation mismatch and parser defects	absence of defects
Statistical experiment	empirical distribution or attack result	binding, zero knowledge, or anonymity as a theorem
Benchmark	performance for a pinned environment/workload	performance on other environments
Code review	implementation findings	cryptographic proof
External audit	independent review within its stated scope	guarantees outside that scope

Miri, sanitizers, linters, tests, and manual review are valuable assurance tools but MUST NOT be called formal verification. One million collision-free trials do not prove commitment binding. A prover test with no observed false acceptance does not establish 2^{128} soundness. Uniform-looking ciphertext bytes do not prove privacy.

4. Requirements traceability

Requirement	Primary tasks	Required release evidence
R1 Private ledger	T103–T104, T301–T305, T401–T405	security argument, vectors, two implementations, audit
R2 Network anonymity	T601–T602, T604	threat model, simulation, adversarial experiments, review
R3 Post-quantum security	T001, T103–T104, T201–T204, T303–T305, T601, T701	composed 128-bit PQ analysis and KATs
R4 Authorization/supply	T202, T301, T305, T402–T405, T504	no-inflation proof/tests and issuance-cap proof
R5 Consensus safety	T403, T501–T504, T603	deterministic models, vectors, reorg/recovery tests

Requirement	Primary tasks	Required release evidence
R6 Scalability	T004, T304–T305, T603–T604	pre-node feasibility and 1,000 accepted tx/s end-to-end artifacts
R7 Setup/upgrade safety	T002, T303–T304, T501, T701	setup analysis, version/replay/upgrade review
R8 Verifiability	T003–T004, all implementation tasks, T701	trace report, independent implementations, external review

A task is not complete if its output is not linked from this matrix through a machine-readable evidence manifest.

5. Acyclic task graph

Tasks MUST be completed in dependency order. A later experiment may send an earlier task back to design, but dependencies MUST NOT be bypassed.

Layer 0 - Scope and evidence

T001 Security profile and threat model	
T002 Canonical serialization and domain registry	<- T001
T003 Test-vector and interoperability harness	<- T002
T004 Benchmark and artifact harness	<- T001, T002

Layer 1 - Algebra and commitment

T101 Field and extension-field arithmetic	<- T001, T003
T102 Polynomial/NTT layer, only if selected	<- T101
T103 Samplers and commitment candidate	<- T001, T002, T101; T102 if required
T104 Commitment security gate	<- T003, T004, T103

Layer 2 - Standardized cryptography and keys

T201 SHAKE256 and domain hashing	<- T002, T003
T202 Hash-based authorization comparison	<- T201
T203 ML-KEM, KDF and authenticated encryption	<- T201
T204 Wallet key hierarchy and address profile	<- T202, T203

Layer 3 - Validity proof

T301 Integer transaction relation and AIR	<- T101, T104, T201, T202
T302 Membership, nullifier and state AIR	<- T301
T303 Fiat-Shamir, FRI and zero-knowledge profile	<- T101, T201
T304 Prover, verifier and aggregation	<- T301, T302, T303

Layer 4 - Transactions and state

T401 Note creation and recipient encryption	<- T104, T203, T204
T305 Cryptographic transaction feasibility gate	<- T004, T202, T204, T304, T401
T402 Transaction codec, builder and authorization hash	<- T202, T301, T305, T401
T403 Commitment/nullifier state and atomic application	<- T302, T402
T404 Wallet scanning and spending	<- T204, T401, T402, T403
T405 Selective disclosure and bounded payment policies	<- T204, T304, T404

Layer 5 - Consensus and issuance

T501 Versioned GHOSTDAG consensus profile	<- T001, T002
T502 Header, PoW and deterministic DAA	<- T201, T501
T503 Canonical DAG-to-state ordering	<- T403, T501, T502
T504 Rewards, supply cap and genesis	<- T503

Layer 6 - Network and integration

T601 Post-quantum authenticated P2P transport	<- T202, T203, T501
T602 Network anonymity protocol	<- T004, T601
T603 Full node, pruning and recovery	<- T305, T503, T504, T601
T604 End-to-end wallet/network/system validation	<- T405, T602, T603

Layer 7 - Independent assurance

T701 Composition review and release decision	<- T104, T305, T504, T604
--	---------------------------

This ordering removes the former cycle in which polynomial multiplication depended on an NTT task that itself depended on polynomial multiplication.

6. Task specifications

T001 — Security profile and threat model

Objective: Freeze the security games, adversaries, corruption thresholds, privacy boundaries, lifetime, multi-user counts, and minimum composed post-quantum security target.

Deliverables:

- cryptographic, consensus, network, side-channel, and operational threat models;
- explicit definition of ledger privacy, recipient-key privacy, and network anonymity;
- per-component and composed security budget;
- quantum collision, preimage, and second-preimage cost models, including time/memory assumptions, multi-target loss, and required digest lengths;
- signed prior-art, reproducibility, reuse-rights, and adopt/adapt/replicate decision for every construction selected for implementation;
- exclusions such as compromised endpoints, with user-facing claim wording.

Pass: independent reviewers agree the model is complete enough to analyze R1–R8. **Reject:** “quantum-safe” or “anonymous” appears without a defined game, adversary, or boundary.

T002 — Canonical serialization and domain registry

Objective: Specify every consensus byte and hash domain before code.

Deliverables:

- field-by-field encodings, endian rules, length widths, maxima, rejection behavior, version and chain binding;
- registry for transaction ID, authorization, note, nullifier, Merkle, transcript, block, PoW, KDF, address, and upgrade domains;
- output length and required security property for every hash domain;
- parser limits and cost ordering;
- replay and downgrade rules.

Pass: a second team can create the same bytes without reading an implementation. **Reject:** implicit serialization, trailing bytes, unbounded vectors, unknown-version fallback, or one hash reused across purposes.

T003 — Test-vector and interoperability harness

Objective: Build the reproducible owner for official KATs and Quantum cross-language vectors.

Deliverables:

- provenance for each external vector;
- canonical JSON or another frozen, integer-safe vector format;
- positive and negative vectors;
- deterministic generator revision and independently checked expected values;
- runner contract for at least two implementations.

Pass: clean-room implementations agree byte-for-byte and reject every negative vector. Generated expected values MUST NOT be produced by the same code path being tested without independent confirmation.

T004 — Benchmark and artifact harness

Objective: Make all performance, storage, and network evidence reproducible.

Deliverables:

- pinned operating systems, compilers, CPU features, memory, disks, network topology, clock source, and revisions;
- deterministic workload generator with private note creation, proofs, conflicts, malformed traffic, delay, and reorgs;
- raw metrics, logs, hashes, manifests, and report generator;
- accepted-state-transition throughput as the primary metric.

Pass: an independent lab reproduces the reported result within a declared tolerance. **Reject:** “latest stable,” unpublished cloud shape, submitted TPS, warm-cache-only measurement, or missing raw artifacts.

T101 — Field and extension-field arithmetic

Objective: Select and implement candidate base and extension fields whose use and security are then approved by T303.

Deliverables: canonical field encoding; add/subtract/multiply/invert; extension arithmetic; batch inversion; constant-time policy; reference and optimized implementations; exhaustive small-model and differential tests.

Pass: all identities, boundary values, invalid encodings, and cross-language vectors pass; proof-system reviewer confirms the field supports the soundness analysis. **Reject:** a 64-bit base field is presented as 128-bit soundness by itself.

T102 — Polynomial and NTT layer

Objective: Implement polynomial and NTT operations only after the selected commitment or proof profile requires them.

Deliverables: ring definition, modulus/root derivation, forward/inverse normalization, naive reference multiplication, optimized multiplication, and aliasing/memory rules.

Pass: optimized results equal the independent naive reference across vectors and randomized boundary cases. **Reject:** circular dependency, floating-point roots, undocumented padding, or silent reduction.

T103 — Samplers and commitment candidate

Objective: Implement one precisely cited commitment candidate without inventing a new formula.

Deliverables: construction citation and delta from it; exact parameters; matrix/seed generation; canonical encodings; unbiased rejection samplers; commit/open API; note-field binding; deterministic test mode isolated from production entropy.

Pass: correctness and vectors pass and the artifact is ready for T104 analysis. **Reject:** byte modulo sampling, reused randomness, invertible-linear “binding” intuition, mismatched randomness length, or carry-dependent value encoding.

T104 — Commitment security gate

Objective: Decide whether the commitment candidate meets R1, R3, and R4.

Deliverables: quantum hiding/binding definitions and reduction; concrete parameter estimates with reproducible lattice-estimator inputs; multi-target loss; side-channel review; independent cryptographer report.

Pass: composed security remains at least 128 post-quantum bits and all findings are resolved.

Reject: collision-search testing is used as proof, the construction is only named “BDLOP-like,” or a guessed value can be tested from public data.

T201 — SHAKE256 and domain hashing

Objective: Implement FIPS 202 SHAKE256 and the Quantum QH wrapper exactly.

Deliverables: official FIPS 202 KATs; incremental/one-shot behavior; Quantum vectors; tag registry enforcement; output-length binding; per-domain collision/preimage/second-preimage requirement; concrete quantum and multi-target output-length analysis; canonical error behavior.

Pass: official and Quantum vectors match in both implementations. **Reject:** SHA3-256 is silently substituted, tags are concatenated ambiguously, or XOF length is unbound or a 32-byte research vector is treated as proof that every collision-dependent consensus digest meets R3.

T202 — Hash-based authorization comparison

Objective: Compare exact FIPS 205 SLH-DSA-SHAKE-256f with an independently specified TzEL-shaped one-time baseline and, if applicable to the wallet model, one exact NIST SP 800-208 stateful profile.

Deliverables: the signed T305 prior-art and reuse decision; exact construction, parameters, source/standard revisions, APIs, contexts, randomization and state-transition choices for every arm; official KATs where available; canonical key/signature encodings; authorization digest; key-erasure and side-channel rules; state-index allocation, crash, backup/restore, concurrency, exhaustion, recovery, and desynchronisation cases; malformed signature/state corpus; per-arm proof-circuit cost report; FIPS 205 message-bound-signature analysis and any reviewed transformation; signed applicability decision for the NIST stateful arm before benchmark results are viewed.

Pass: both implementations match applicable KATs and shared vectors, reject malformed/non-canonical inputs and state misuse, and expose all retained arms to T305 under the same transaction relation. **Reject:** protocol IDs say SPHINCS+, TzEL code is copied without compatible permission and legal review, a classical fallback exists, signatures authorize a partial transaction, or stateful failure modes are omitted.

T203 — ML-KEM, KDF and authenticated encryption

Objective: Define recipient encryption building blocks with ML-KEM-1024.

Deliverables: FIPS 203 KATs and errata version; decapsulation-failure behavior; transcript-bound KDF; 256-bit AEAD profile; nonce and replay rules; key separation; receiver-anonymity/key-privacy game for note encryption; multi-user, chosen-key, chosen-ciphertext and cross-output analysis; scanning-tag disclosure budget; test vectors.

Pass: chosen-ciphertext, receiver-anonymity and misuse analysis covers the exact note-encryption composition. **Reject:** ML-KEM is treated as authentication, failure creates an oracle, nonce reuse is possible, or an invented Noise pattern name substitutes for a message specification, or the ciphertext identifies its recipient key.

T204 — Wallet key hierarchy and address profile

Objective: Freeze mnemonic import, child-key derivation, address payload, and storage protection.

Deliverables: exact BIP-39 behavior; domain-separated post-quantum child derivation; spend/nullifier/encryption key separation; incoming-view, full-wallet-view, transaction-disclosure,

and auditor-scoped capability derivation; canonical scope encodings; canonical lower Base32 address vectors; memory-hard wallet-storage KDF; recovery tests.

Pass: BIP-39 official vectors and Quantum derivation/address/disclosure vectors match both implementations, and each disclosure capability reveals no more than its stated scope. **Reject:** PBKDF2 parameters are called BIP-39 while using a different salt/normalization, one key crosses protocol roles, or a universal viewing backdoor is required.

T301 — Integer transaction relation and AIR

Objective: Encode authorization, input/output openings, encryption binding, ranges, and exact conservation.

Deliverables:

- formal relation before AIR translation;
- public-input and witness schemas;
- constraints for every input commitment opening and membership leaf;
- constraints for every output commitment opening;
- complete verification of each retained T202 authorisation profile over the same complete digest;
- 16-bit limb range checks and wide carry-constrained sums;
- fee and reward-type separation;
- AIR degree and trace-size report.

Required negative cases:

- replace an input value while keeping its commitment;
- spend a one-unit note while claiming a near-maximum input value;
- make both sides equal only modulo the proof field;
- cross a 16-bit carry boundary such as $65,535 + 1$;
- omit or alter an output opening, fee, chain ID, or authorization field.

Pass: every negative case is unsatisfiable and reviewers trace each R4 clause to constraints.

Reject: balance is only a field equality or an input opening is absent.

T302 — Membership, nullifier and state AIR

Objective: Prove note membership and correct unlinkable nullifier derivation against the exact state hashes.

Deliverables: depth and empty-root definitions; leaf/internal tags; membership-index constraints; nullifier relation; local uniqueness; anchor/public-state binding; digest length derived from T001/T201 for Merkle collision and second-preimage requirements; invalid-path and alternate-tree vectors.

Pass: all paths and indices are bound, invalid paths fail, and no public field creates a direct note-to-nullifier link. The selected digest length MUST meet R3 after multi-target and lifetime losses. **Reject:** path direction is unconstrained, empty roots differ across implementations, a 32-byte wrapper vector silently becomes the Merkle security parameter, or tree depth exhaustion is ignored.

T303 — Fiat-Shamir, FRI and zero-knowledge profile

Objective: Select the complete transparent proof profile.

Deliverables: base/extension fields; transcript schedule; challenge sampling; AIR composition; FRI/DEEP-FRI parameters; query/grinding/batching; zero-knowledge masks; classical/QROM soundness calculations; multi-proof loss; denial-of-service limits.

Pass: independent proof-system review shows at least 128 post-quantum bits after composition and no witness leakage under the stated model. **Reject:** the arithmetic “queries $\times$ bits” is used as soundness, target is 2^{100} , masking is absent, or a generic STARK is assumed zero knowledge.

T304 — Prover, verifier and aggregation

Objective: Implement the T301–T303 profile without accepting ambiguous proofs.

Deliverables: independent prover/verifier interfaces; streaming/resource limits; transcript vectors; malformed-proof corpus; deterministic failure codes; individual and, if selected, aggregate wire mode; exact binding of an aggregate to the ordered transactions, public inputs, and pre/post-state; differential/fuzz tests; measured proof generation and verification.

Pass: invalid witnesses and mutated proofs fail; both implementations interoperate; aggregate mode omits replaced individual proofs; performance is fed into T305. **Reject:** verifier panics, work is allocated before size checks, both proof modes are redundantly transmitted, or tests are claimed as soundness proof.

T305 — Cryptographic transaction feasibility gate

Objective: Decide whether the complete private transaction is practical before wallet, node, or GHOSTDAG integration consumes implementation effort.

Experiment protocol: The non-normative T305 research protocol records the prior-work boundary, pre-registration, independent-implementation method, evidence package, and stop/go form. It may strengthen experimental detail but cannot weaken this task or the protocol specification.

Required prototype: exactly two inputs and two outputs with input openings, Merkle membership, nullifier derivation, output openings, encrypted-note binding, complete in-proof authorisation for each retained T202 arm, 64-bit range constraints, carry-safe conservation, public fee, zero-knowledge masking, and the selected individual and aggregate proof path. Only authorisation and its required state may differ between arms.

Deliverables: two interoperable implementations; pinned desktop and constrained-client hardware; single-wallet p50/p95/p99 proving latency; peak and steady memory; verifier time and memory; proof and transaction bytes; parallel-prover scaling; aggregate latency, size, and amortised verification; isolated authorisation contribution; state-management failure results; pre-registered material-benefit criteria; clearly labelled reproduced versus author-reported prior-work measurements; raw artifacts; and a signed GO/ADAPT/REPLICATE/STOP decision.

Pass: every T004 threshold frozen before the run is met without mock signatures, omitted relations, disabled zero knowledge, retained individual proofs in aggregate mode, unreported preprocessing, or post-result benefit criteria. Retaining the stateless profile additionally requires a material pre-registered benefit over every qualifying stateful arm. **Reject:** no arm meets the frozen client, verifier, wire, and aggregate budgets. A reject returns the signature, commitment, relation, proof profile, or explicit system requirements to research; it MUST NOT be deferred to node optimisation.

T401 — Note creation and recipient encryption

Objective: Create a note that the recipient can discover and later spend without publishing recipient or amount.

Deliverables: note plaintext schema; commitment randomness; encryption witness; diversified recipient data; authenticated binding to transaction and output index; receiver-anonymity/key-privacy analysis linked to T001/T203; sender/recipient test vectors; failure and deletion behavior.

An output-creation API MUST receive all data needed by the final note commitment, including spend-authorization and nullifier-key digests. A view public key alone is insufficient.

Pass: intended recipients recover exactly the committed note, non-recipients cannot distinguish which eligible recipient key was targeted beyond the approved disclosure budget, and replay/swapping fails. **Reject:** encryption and commitment can describe different notes, or a ciphertext/scanning tag identifies the recipient.

T402 — Transaction codec, builder and authorization hash

Objective: Construct and parse the exact public transaction while keeping private witnesses off-ledger.

Deliverables: canonical codec; maximum counts/sizes; dedicated sighash; builder state machine; duplicate handling; expiry/anchor/authorisation-profile/proof-mode binding; negative parser corpus.

Pass: round trips are byte-stable, alternate encodings fail, and every mutable semantic field changes the authorization digest. **Reject:** parser fallbacks, unbounded allocation, signature malleability, or unknown fields silently ignored.

T403 — Commitment/nullifier state and atomic application

Objective: Implement the single owner for note-tree, nullifier-set, and state-root transitions.

Deliverables: atomic apply/revert; deterministic roots; duplicate nullifier checks; snap-shot/recovery format; reorg journal; concurrency model; invariants and crash tests.

Transaction verification MUST take an explicit immutable state/anchor context. A stateless `Transaction::verify()` is insufficient for global nullifier and anchor checks.

Pass: crashes and reorgs recover to the same root, two spends cannot both commit, and no non-owner bypasses state validation. **Reject:** check-then-write races, partial state commits, or hidden empty-state fallback.

T404 — Wallet scanning and spending

Objective: Build deterministic private wallet behavior over T401–T403.

Deliverables: scan/decrypt pipeline; note and nullifier tracking; spend selection; proof/sign/build flow; reorg rollback; scoped disclosure export and receipt verification; disclosure limitations; secure deletion; user-facing finality and fee semantics.

Pass: restore-from-seed and rescan reproduce the same wallet state through reorgs; spent notes cannot be selected; no network identity is derived from wallet keys. **Reject:** missing data is zero-filled, failed decryption is treated as an empty note, or sensitive logs are emitted.

T405 — Selective disclosure and bounded payment policies

Objective: Add useful settlement controls without introducing a general virtual machine, a universal viewing key, or an unbounded proof relation.

Deliverables: incoming-view, full-wallet-view, transaction-specific, and auditor-scoped disclosure flows; canonical scope and receipt encodings; least-privilege and unlinkability tests; documented non-revocation of copied history; a finite registry of accepted payment policies; exact semantics and costs for each enabled timelock, hashlock, threshold/multisignature, escrow, recurring-authorisation, atomic-swap, or conditional-release policy; updated AIR/proof profile, vectors, and resource report for every enabled policy.

Private asset issuance and cross-chain adapters are separate post-core profiles. Any proposed adapter MUST publish its external consensus, custody, MPC/multisignature, oracle, light-client, finality, privacy, and upgrade trust assumptions; it does not inherit Quantum’s base-layer assurance.

Pass: disclosure never exceeds its canonical scope, policy execution is deterministic and bounded across reorgs/timeouts, and every enabled policy preserves transaction privacy, authorisation, supply, and T305 performance budgets. **Reject:** a global backdoor key exists, copied disclosure is called revocable, unknown policy bytecode falls back to execution, a policy can mint or bypass authorisation, or an adapter is described with stronger properties than its own evidence supports.

T501 — Versioned GHOSTDAG consensus profile

Objective: Turn the cited GHOSTDAG research into one deterministic, versioned protocol profile.

Deliverables: parent validation; anticone and blue/red rules; selected parent; merge ordering; blue work; finality/pruning assumptions; block limits; reference model; exhaustive small-DAG and adversarial vectors.

Pass: independent implementations produce identical classification, ordering, score/work, and finality across all vectors. **Reject:** “highest blue score wins,” a simple topological sort, or upstream behavior without a pinned version is treated as the specification.

T502 — Header, proof of work and deterministic DAA

Objective: Specify and implement canonical headers, PoW, and expected difficulty.

Deliverables: exact header bytes and size; block/PoW domains; integer target conversion; DAA equations, window, clamps and overflow rules; timestamp rules; genesis-independent vectors; time-warp simulations.

Required negative cases:

- miner supplies an easy target;
- nodes have different local clocks;
- timestamps are reordered or clamped;
- maximum/minimum targets and arithmetic boundaries;
- semantically equal headers use alternate bytes.

Pass: validators recompute and enforce the same expected target without floating point or local-time dependence. **Reject:** the header target is only used to check its own PoW.

T503 — Canonical DAG-to-state ordering

Objective: Prove and implement one deterministic state transition order for parallel blocks.

Deliverables: ordered merge-set algorithm; transaction ordering; duplicate and concurrent-nullifier rule; pre/post-state commitments; reorg semantics; model checker or equivalent state exploration; adversarial DAG vectors.

Pass: all implementations converge on accepted transactions and state root under permutations, delays, conflicts, and reorgs. **Reject:** parallel blocks mutate shared state concurrently or conflict resolution depends on arrival order.

T504 — Rewards, supply cap and genesis

Objective: Define value creation and the immutable network start.

Deliverables: distinct reward transaction; DAA-score or other exact DAG issuance index; blue/red/stale eligibility; integer schedule and rounding; cumulative cap proof; exact fee-transfer, subsidy-issuance, burn and note-supply equations; reward maturity and reorg reversal; private reward-output handling; canonical genesis bytes and vectors.

Pass: exhaustive/bounded-model checks and analytical proof show issuance never exceeds 21,000,000 QTM across reorganizations, fees are never counted as new issuance, and all nodes derive identical rewards. **Reject:** linear height is used without a DAG definition, a fee can be claimed twice or counted as subsidy, no coinbase/reward transaction exists, timestamp units differ, or genesis hashes with another function.

T601 — Post-quantum authenticated P2P transport

Objective: Define link confidentiality/authentication without classical fallback or wallet identity leakage.

Deliverables: exact handshake message sequence; ML-KEM/SLH-DSA transcript binding; KDF/AEAD keys; replay/downgrade/key-compromise analysis; peer identity lifecycle; vectors and state-machine fuzzing.

Pass: independent cryptographic review approves the exact composition and implementations interoperate. **Reject:** “post-quantum Noise” is used without an actual reviewed pattern, ML-KEM is mistaken for authentication, or failed PQ negotiation falls back to X25519.

T602 — Network anonymity protocol

Objective: Meet R2 against the T001 observer rather than merely encrypt links.

Deliverables: routing and cover-traffic protocol; peer/route selection; batch/delay distributions; active tagging, eclipse, Sybil and intersection analysis; simulator; reproducible attack experiments; utility/privacy tradeoff report.

Metrics: source-classification advantage, anonymity-set distribution, entropy, precision/recall for origin inference, latency, bandwidth amplification, failure behavior, and confidence intervals.

Pass: thresholds defined in T001 are met under all required scenarios and an independent network-privacy review passes. **Reject:** only ciphertext byte uniformity, payload correlation, or Dandelion++ deployment is offered as proof of full anonymity.

T603 — Full node, pruning and recovery

Objective: Integrate consensus, proof verification, atomic state, transport, storage, and resource admission.

Deliverables: validation pipeline with cheap checks first; bounded queues; peer resource accounting; state persistence; pruning/snapshot/archive model; trustless recovery; reorg handling; observability without sensitive data; fault-injection results.

Pass: state roots remain identical under crash, restart, overload, reorder, and recovery; no malformed input causes unbounded work. **Reject:** hidden fallback state, partial validation, archive size presented as pruned-node size, or snapshot trust is undocumented.

T604 — End-to-end wallet/network/system validation

Objective: Demonstrate R1–R8 together, not in isolated microbenchmarks.

Required run:

- geographically distributed topology;
- private note creation, scanning, spending, proofs, propagation and state application enabled;
- selective-disclosure flows and every payment policy enabled in the tested release profile;
- at least 1,000 accepted layer-1 tx/s for 24 hours;
- conflict, delay, peer loss, malformed traffic, reorg, restart and recovery phases;
- raw proof-size, bandwidth, storage, CPU, memory, latency, rejection and anonymity metrics.

Pass: all functional invariants remain true, the throughput target is sustained, the anonymity thresholds pass, and another team reproduces the report. **Reject:** benchmark uses mock proofs, prevalidated transactions, disabled privacy, submitted TPS, a single process, or unpublished hardware.

T701 — Composition review and release decision

Objective: Decide whether the combined system—not only its parts—meets the three non-negotiable requirements.

Deliverables: requirements-to-evidence manifest; independent cryptographic, proof, consensus, network, implementation and performance reports; finding register; per-capability claim register; residual-risk statement; legal review; signed go/no-go record.

Pass for testnet: all blocking design gates pass, two implementations interoperate, and no unresolved critical/high issue permits inflation, unauthorized spend, privacy break, consensus divergence, or target bypass.

Pass for production: testnet evidence remains valid over the approved stability period, external audits are closed, governance/upgrades are post-quantum and replay-safe, operational recovery is proven, and public claims match evidence.

Reject: an author or AI self-certifies the system, a component review is presented as composition review, or any hard requirement is deferred until after launch.

7. Mandatory adversarial scenarios

Every relevant implementation must preserve a permanent regression test for:

ID	Scenario	Required outcome
A001	Small real input opened with a larger private input value	Proof unsatisfiable
A002	Input/output balance equal only modulo the proof field	Proof unsatisfiable
A003	65,535 + 1 crosses an unconstrained limb carry	Correct integer result only
A004	Output commitment and encrypted output describe different notes	Rejected
A005	Authorization omits fee, output, chain, anchor, authorisation profile, or proof mode	Rejected
A006	Same nullifier in one transaction, block, parallel block, or reorg	At most canonical first spend accepted
A007	Header declares an easier target than the DAA computes	Block rejected
A008	Honest nodes use different local clocks on the same DAG	Identical consensus result
A009	Equivalent DAG delivered in different arrival orders	Identical ordering and state root
A010	Oversized vectors/proofs or unknown version	Rejected before expensive work
A011	Failed PQ negotiation requests classical fallback	Connection rejected
A012	Origin observer combines timing, volume and route metadata	Must remain below T001 attack threshold
A013	Snapshot interrupted or maliciously altered	Recovery rejects or reaches canonical root
A014	Reward rounding/reorg approaches the supply cap or replays fees	Cap never exceeded; fee portion never counted as issuance or claimed twice
A015	Observer compares note ciphertexts or scanning tags against candidate recipient keys	Recipient inference remains below the T001 threshold
A016	Disclosure capability is replayed outside its account, counterparty, transaction, or time scope	No additional note or wallet history is revealed
A017	Unknown payment-policy identifier, unbounded policy input, timeout/reorg edge, or adapter downgrade	Fail closed; deterministic bounded state; no inherited bridge claim

These are minimum cases, not a complete security test suite.

8. Performance and capacity report

Every benchmark report MUST contain:

1. repository revisions and build hashes;
2. exact dependency/toolchain versions;
3. machine and network inventory;
4. workload seed and transaction-shape distribution;
5. accepted, rejected, duplicate, and submitted counts separately;
6. p50/p95/p99 proof, propagation, ordering, application and finality latency;
7. average/p95/max transaction and proof bytes;
8. per-node inbound/outbound bandwidth and gossip amplification;
9. CPU, memory, disk I/O, state and archive growth;
10. anonymity metrics and added bandwidth/latency;
11. all failures, restarts, divergences and discarded samples;
12. commands and immutable artifact locations.

T305 reports MUST additionally separate single-wallet latency from parallel prover throughput and aggregate amortisation, report authorisation cost per arm, and label prior-work numbers as reproduced or author-reported. A server aggregate rate MUST NOT be reported as client-side transaction latency.

The 1,000 tx/s acceptance target is fixed. Reference hardware and proof-latency thresholds are intentionally not fabricated in advance; T004 must freeze them before a result can be interpreted. A 60-second benchmark limit MUST NOT be quietly substituted for a different specification limit.

9. Security review checklist

Cryptographic correctness

- ☐ FIPS 202, 203, and 205 algorithm names and parameter sets are exact.
- ☐ Official KAT provenance is recorded.
- ☐ Every hash domain has a security property and output length justified against quantum and multi-target attacks.
- ☐ Commitment construction and parameters pass T104.
- ☐ Every random sampler is unbiased and domain-separated.
- ☐ No secret-dependent branch, memory lookup, log, or error oracle remains.
- ☐ Multi-user and protocol-lifetime composition remains 128 PQ bits.

Transaction soundness

- ☐ Each input value is bound to an opened member commitment.
- ☐ Each output value is bound to its public output commitment.
- ☐ Authorization key is bound to the spent note.
- ☐ Authorization digest covers every semantic field.
- ☐ All ranges and sums use constrained integer limbs and carries.
- ☐ Reward minting is type-separated from ordinary conservation.
- ☐ Global anchor/nullifier checks use explicit state context.

Privacy and anonymity

- ☐ No transparent transfer mode exists.
- ☐ Public ledger fields meet the R1 disclosure budget.
- ☐ Note encryption and scanning metadata meet the recipient-key-privacy game.

- ☐ Note ciphertext and commitment cannot be swapped or replayed.
- ☐ Wallet, transport, and transaction identities are separated.
- ☐ T602 covers global passive and required active attacks.
- ☐ Incoming, full-wallet, transaction-specific, and auditor-scoped disclosure are distinct and least-privilege.
- ☐ Copied disclosure is not described as revocable or as proof of legal compliance.

Product boundary and performance

- ☐ T305 passes before full-node or GHOSTDAG integration starts.
- ☐ The 2-input/2-output prototype includes every required relation and each retained in-proof authorisation arm.
- ☐ Stateless SLH-DSA is retained only if its frozen material-benefit rule passes against every qualifying stateful comparator.
- ☐ Stateful key reuse, exhaustion, crash, restore, rollback, concurrency, recovery, and desynchronisation cases are tested.
- ☐ Wallet latency, parallel proving, verifier cost, wire size, and aggregate amortisation are reported separately.
- ☐ Every enabled payment policy has bounded deterministic semantics and an updated proof/resource report.
- ☐ No general-purpose VM, private asset, bridge, or adapter inherits a base security, privacy, or compliance claim.

Consensus and supply

- ☐ GHOSTDAG profile is executable and versioned.
- ☐ DAA is deterministic integer arithmetic.
- ☐ Validators recompute expected target.
- ☐ DAG ordering yields one atomic state history.
- ☐ Reward eligibility and index are DAG-native.
- ☐ Fee transfer, subsidy issuance, burns and reward maturity are exact.
- ☐ Cap proof includes fee replay, rounding and reorganization.
- ☐ Genesis bytes and timestamp units are canonical.

Parsing and operations

- ☐ All lengths and counts are bounded.
- ☐ Unknown versions and non-canonical encodings fail closed.
- ☐ Cheap validation precedes expensive cryptography.
- ☐ Queues, memory, CPU and disk work are bounded per peer.
- ☐ Snapshots are authenticated and trustless recovery is tested.
- ☐ Logs and crash artifacts contain no wallet or network-origin secrets.

Unchecked boxes are normal at this research stage. They MUST NOT be prechecked in reports.

10. Evidence manifest and report format

Each task produces a record with:

```
task_id: T301
status: NOT_STARTED # NOT_STARTED | IN_PROGRESS | BLOCKED | PASS | FAIL
owner: null
reviewer: null
requirements: [R1, R3, R4]
source_revision: null
toolchain_manifest: null
evidence:
```

```

analytical_argument: null
formal_artifact: null
test_report: null
vector_set: null
benchmark_artifacts: null
external_review: null
open_findings: []
decision_date: null
decision_signature: null

```

Null fields mean **not supplied**, never “not needed” or “passed.” PASS is allowed only when every mandatory artifact for that task exists and the named reviewer signs the exact revision.

Current aggregate report

Area	Status	Evidence
Requirements draft	IN_PROGRESS	Specification 0.4.0-research
Commitment	NOT_STARTED	No selected construction
Proof profile	NOT_STARTED	No frozen parameters or AIR
Private transaction	NOT_STARTED	Relation drafted; no implementation
DAG/state integration	NOT_STARTED	Requirements only
Network anonymity	NOT_STARTED	Requirements only
1,000 accepted tx/s	NOT_STARTED	No benchmark
Independent implementations	NOT_STARTED	None
External audit	NOT_STARTED	None

11. Proposed implementation workspace

The following layout is a proposal for a future implementation repository; it does not exist here and is not a command contract:

```

quantum-protocol/
  spec/                canonical encodings and generated vectors
  reference/           simple independent reference implementation
  node/               production-oriented validator implementation
  wallet/             wallet, scanner and proof client
  formal/             models and formal artifacts
  tests/
    kats/
    interoperability/
    adversarial/
    fuzz/
    performance/
  evidence/           immutable manifests and signed reports

```

Build, test, fuzz, benchmark, and vector commands MUST be defined by that repository and pinned in T003/T004. This documentation repository intentionally does not publish fictitious commands that appear runnable.

12. Release labels

Only these labels may be used:

Label	Minimum meaning
Research design	Requirements and candidates; blocking gates may be open
Prototype	Non-production code; no security or performance claim
Reviewed testnet candidate	T701 testnet pass on an exact revision
Public testnet	Candidate deployed with explicit residual risks
Production candidate	T701 production evidence complete; no live claim yet
Production	Exact audited revision deployed and continuously monitored

At the current revision the project is a **research design**.

13. Revision record

0.4.0-research — 2026-07-23

- required a signed revision of the T305 prior-art and reuse decision as a T001/T202 deliverable;
- reframed T202 and T305 as a same-relation comparison between the FIPS 205 stateless incumbent, an independently specified TzEL-shaped baseline, and an applicable NIST SP 800-208 stateful profile;
- required state-management failure vectors, isolated authorisation costs, and reproduced-versus-author-reported evidence labels;
- introduced a requirement to pre-register the material-benefit rule and explicit GO/ADAPT/REPLICATE/STOP outcome classes.

0.3.0-research — 2026-07-21

- expanded the acyclic plan to 31 tasks with T305 as the mandatory complete transaction feasibility gate before node/consensus integration;
- separated single-wallet proof latency, parallel prover throughput, aggregate amortisation, and verifier/wire budgets;
- added T405 for least-privilege disclosure and a finite, proof-bounded payment policy registry;
- made private assets and interoperability adapters separate post-core profiles with explicit external trust and claim boundaries;
- added adversarial cases for disclosure-scope escalation and unknown or unbounded payment-policy behavior.

0.2.0-research — 2026-07-09

- changed the guide from autonomous AI implementation instructions to a human-owned research and verification plan;
- defined all 29 task IDs and removed missing-task references;
- made the dependency graph acyclic;
- added requirements traceability and honest NOT_STARTED reporting;
- added direct regression cases for input-value inflation, field wrap, self-declared PoW targets, local-time DAA divergence, DAG conflicts, and anonymity correlation;
- distinguished proof, formal verification, testing, statistics, benchmarks, review, and audit;
- aligned performance acceptance with 1,000 end-to-end accepted tx/s;
- required external specialist and composition review before any release claim.
- added explicit gates for quantum digest lengths and encrypted-note recipient-key privacy;
- defined fee transfer, claimed subsidy, burns, maturity, and reorg accounting as mandatory supply evidence.